

A Better Lower Bound for On-Line Scheduling

Yair Bartal ^{*} Howard Karloff [†] Yuval Rabani [‡]

Abstract

We consider the on-line version of the original m -machine scheduling problem: given m machines and n positive real jobs, schedule the n jobs on the m machines so as to minimize the makespan, the completion time of the last job. In the on-line version, as soon as a job arrives, it must be assigned immediately to one of the m machines.

We study the competitive ratio of the best algorithm for m -machine scheduling. The largest prior lower bound was that if $m \geq 4$, then every algorithm has a competitive ratio at least $1 + 1/\sqrt{2} \approx 1.707$. We show that if m is large enough, the competitive ratio of every algorithm exceeds 1.837. The best upper bound on the competitive ratio is now 1.945.

1 Introduction

The m -machine scheduling problem is one of the most widely-studied problems in computer science ([4, 6, 7, 8] are surveys), with an almost limitless number of variants. In this note, we study one of the simplest and earliest m -machine scheduling problems ever studied, the scheduling problem of Graham, introduced in 1966 [3]. This is the variant in which each job consists of exactly one task, which requires the same execution time on each of the m machines. Jobs cannot be preempted and are independent of each other. The goal is to minimize the *makespan*, the completion time of the last job. Formally, the problem is this: Given a sequence of positive reals a_1, a_2, \dots, a_n and an integer m , for each j assign a_j to a machine i , $1 \leq i \leq m$, so as to minimize the maximum, over i , of the sum of all reals assigned to machine i . Even the special case of $m = 2$ is NP-Hard, as it is at least as hard as PARTITION.

We study the on-line version of the problem: at the time when job j is scheduled, the scheduling algorithm knows only a_1, a_2, \dots, a_j . As soon as a_j appears, the scheduling algorithm,

^{*}Computer Science Department, School of Mathematics, Tel-Aviv University, Tel-Aviv 69978, Israel.

[†]College of Computing, Georgia Institute of Technology, Atlanta, Georgia 30332-0280. The author was supported in part by NSF grant CCR 9107349.

[‡]International Computer Science Institute, 1947 Center Street, Berkeley, CA 94704-1105. Part of this work was done while the author was a student at Tel-Aviv University. The work at ICSI was supported in part by a Rothschild postdoctoral fellowship.

with only the knowledge of a_1, a_2, \dots, a_j , must irrevocably assign job j to one of the machines.

For a sequence σ of jobs, let $A(\sigma)$ denote the makespan of the schedule generated by algorithm A , and let $OPT(\sigma)$ denote the minimum makespan among all m -machine schedules for σ . A 's *competitive ratio* is then

$$c_A := \sup_{\sigma} \frac{E[A(\sigma)]}{OPT(\sigma)},$$

where the supremum is over all nonempty sequences of jobs.

Graham showed in [3] that the naive List algorithm—always place the next job on the most lightly loaded machine—has competitive ratio $2 - 1/m$ for each m . Until recently, no algorithm was known whose performance ratio, as a function of m , was bounded above by $2 - \epsilon$ for all m (for a fixed positive ϵ). Bartal, Fiat, Karloff and Vohra recently exhibited an on-line algorithm with competitive ratio at most $2 - 1/70$ for all m [1]. Currently the best algorithm for large m is due to Karger, Phillips and Torng [5], and has competitive ratio bounded by 1.945.

But how *small* can c_A be? Prior to this note, the best known lower bounds were due to Faigle, Kern and Turan [2], who proved in 1989 that no on-line algorithm could have a competitive ratio that is smaller than $3/2$ if $m = 2$ or smaller than $5/3$ if $m = 3$, and that for no $m \geq 4$ could the competitive ratio be less than $1 + \frac{1}{\sqrt{2}} = 1.707\dots$. This last fact was proven by examining the job sequence consisting of m jobs of size 1, m jobs of size $1 + \sqrt{2}$, and then one job of size $2 + 2\sqrt{2}$.

We improve the lower bound due to Faigle, Kern and Turan to $\alpha \approx 1.837$. The proof is quite simple: we give the algorithm one sequence of judiciously-chosen jobs, and then argue that on at least one nonempty prefix of the sequence, the makespan obtained by algorithm A divided by the optimal makespan for those jobs is at least α . Since the algorithm is on-line, its schedule for jobs $\sigma_1, \sigma_2, \dots, \sigma_i$ is obtained by adding job σ_i to its schedule for jobs $\sigma_1, \sigma_2, \dots, \sigma_{i-1}$.

2 The Construction

We will give the on-line algorithm a sequence of jobs, whose sizes depend on the positive reals x, y, z , whose values we will choose later. Rather than give all the constraints that must be satisfied by x, y , and z now, we give them “on-the-fly” when they are needed. These parameters are chosen so that any algorithm that has competitive ratio less than α must schedule the jobs in a unique way. In fact, it must schedule each new job on the lowest machine available.

The job sequence is as follows.

1. m jobs of size $1/(x + 1)$.
2. m jobs of size $x/(x + 1)$.

3. m jobs of size x .
4. $\lfloor m/2 \rfloor$ jobs of size y .
5. $\lfloor m/3 \rfloor - 2$ jobs of size z .
6. $m + 3 - \lfloor m/2 \rfloor - \lfloor m/3 \rfloor$ jobs of size $2y$.

Now we study the performance of deterministic on-line algorithm A . We assume, for a contradiction, that its competitive ratio is less than α (where α is as yet undetermined). Let us use On_i and Opt_i to denote the on-line and optimal makespan, respectively, for the job sequence consisting of blocks $1, 2, \dots, i$, $1 \leq i \leq 6$.

If the m jobs of size $1/(x+1)$ are not placed on separate machines, then $On_1 \geq 2/(x+1)$. Clearly $Opt_1 = 1/(x+1)$, so that in this case the competitive ratio is at least α if $\alpha \leq 2$. So we stipulate that

$$\alpha \leq 2, \tag{1}$$

thereby forcing A to place the m $1/(x+1)$'s on separate machines.

Now if the m $x/(x+1)$'s aren't on separate machines, then $On_2 \geq 2x/(x+1) + 1/(x+1) = (2x+1)/(x+1)$, and clearly $Opt_2 = 1$. So let us stipulate that

$$\frac{2x+1}{x+1} \geq \alpha, \tag{2}$$

to force $On_2 = 1$.

We do the same with the x 's. Equation (2) fortuitously forces A to place the x 's on separate machines. This means that $On_3 = 1+x$; clearly $Opt_3 = 1+x$ also.

Now we have to force A to place the $\lfloor m/2 \rfloor$ y 's on separate machines. What is Opt_4 ? After the y 's arrive, we can place one $1/(x+1)$ and one $x/(x+1)$ on each machine, for a total height of 1. Next, we can place the x 's on $\lceil m/2 \rceil$ machines, at most two x 's per machine, so that those machines have height $1+2x$. On the remaining $m - \lceil m/2 \rceil = \lfloor m/2 \rfloor$ machines, we place one y per machine. This means that $Opt_4 \leq \max\{2x+1, y+1\}$. We will want to use $Opt_4 \leq 2x+1$, so we add the condition

$$y \leq 2x. \tag{3}$$

If A places any two of the y 's atop one another, then $On_4 \geq 2y+x+1$. However, $Opt_4 \leq 2x+1$. By stipulating that

$$\frac{2y+x+1}{2x+1} \geq \alpha, \tag{4}$$

we force A to place the y 's on separate machines.

Now the z 's, of which we have $\lfloor m/3 \rfloor - 2$. We want A to place them alongside each other, atop the x 's. If it doesn't, then its makespan will be at least $1+x+y+z$. We will ensure that

$Opt_5 \leq 3x$. Once this is done, we will stipulate that

$$\frac{z + y + x + 1}{3x} \geq \alpha, \quad (5)$$

to force A to place the z 's alongside each other, atop the x 's.

One way to pack the items in blocks 1, 2, 3, 4, 5 is as follows. $\lceil m/6 \rceil$ machines have three x 's each. $\lfloor m/2 \rfloor$ machines have a y , an x , an $x/(x+1)$, and two $1/(x+1)$'s. $\lfloor m/3 \rfloor - 2$ machines have a z and two $x/(x+1)$'s. The reader can verify that no more than m machines have been used, and that all jobs in blocks 1-5 have been scheduled, provided that $m \geq 32$. The makespan of this schedule is $\max\{z + 2x/(x+1), y + x + x/(x+1) + 2/(x+1), 3x\}$. We want this to be $3x$. So we stipulate that

$$z + 2\frac{x}{x+1} \leq 3x \quad (6)$$

and

$$y + x + 1 + \frac{1}{x+1} \leq 3x. \quad (7)$$

Now, the last block, the $(2y)$'s. Notice that the number of y 's, z 's and $(2y)$'s, summed up, is $m + 1$. This means that A must place at least one of the $(2y)$'s above either a y or a z . In any case, $On_6 \geq 3y + x + 1$, if we stipulate that

$$y \leq z. \quad (8)$$

We will add enough constraints to ensure that $Opt_6 = 2y$. So we will have a contradiction if we stipulate that

$$\frac{3y + x + 1}{2y} \geq \alpha. \quad (9)$$

Here is a partial schedule at the end: There are $m + 3 - \lfloor m/2 \rfloor - \lfloor m/3 \rfloor$ $(2y)$'s, which go on separate machines. The $\lfloor m/2 \rfloor$ y 's go on $\lceil m/4 \rceil$ machines, with up to two y 's apiece. $\lceil m/3 \rceil$ machines have one z and one x . $\lceil 2m/9 \rceil$ machines have three x 's and three $x/(x+1)$'s apiece.

On the remaining machines (which number at least $-8 + m/36$), we must schedule m $1/(x+1)$'s and $m - 3\lceil \frac{2m}{9} \rceil$ jobs of size $x/(x+1)$. On $\lceil m/63 \rceil$ machines, we place 21 $x/(x+1)$'s apiece. On $\lceil m/111 \rceil$ machines we place 111 $1/(x+1)$'s. This is a valid schedule, provided that we have $\lceil m/63 \rceil + \lceil m/111 \rceil \leq -8 + m/36$. This is certainly true if $m \geq 3454$.

The makespan of this schedule is

$$\max\{2y, z + x, 3x + 3\frac{x}{x+1}, 21\frac{x}{x+1}, 111\frac{1}{x+1}\}.$$

Since we want the makespan to be $2y$, we stipulate that

$$z + x \leq 2y, \quad (10)$$

$$3x + 3\frac{x}{x+1} \leq 2y, \quad (11)$$

$$21\frac{x}{x+1} \leq 2y, \quad (12)$$

$$111\frac{1}{x+1} \leq 2y. \quad (13)$$

What remains is to find x, y, z, α satisfying all the inequalities above, with α as large as possible. We will arbitrarily choose a set of inequalities to make tight, solve them, and show numerically that the remaining constraints are satisfied. We choose inequalities (2), (5), (9), and (10) to be satisfied with equality; the others will not be. From equations (2), (5), and (10), we get

$$\frac{2x+1}{x+1} = \alpha = \frac{3y+1}{3x},$$

which yields

$$y = \frac{1}{3} \left(\frac{6x^2 + 2x - 1}{x+1} \right). \quad (14)$$

From equations (5), (9), and (10) we get

$$\frac{3y+1}{3x} = \frac{3y+x+1}{2y}. \quad (15)$$

Eliminating y from equations (14) and (15) we get

$$3x^3 - 13x^2 - 12x - 2 = 0.$$

It is not hard to verify that there are roots in the intervals $(-0.58023953020, -0.58023953019)$, $(-0.2236520327, -0.2236520326)$, and $(5.137224896, 5.137224897)$. We choose x to be the root in the final interval. In fact, the exact value for x is

$$\frac{13}{9} + \frac{2}{9}\sqrt{277} \cos \left(\frac{1}{3} \arccos \frac{4546}{277\sqrt{277}} \right).$$

This yields $y \approx 9.104056552$ and $z \approx 13.070888209$. This gives us $\alpha = (2x+1)/(x+1) = 2 - 1/(x+1) \in (1.8370599062, 1.8370599063)$. It is now easy to verify numerically that the remaining constraints are all strictly satisfied. This gives us a lower bound exceeding 1.837.

3 Acknowledgments

We are grateful to Steven Phillips and Eric Torng for improving our earlier results.

References

- [1] Y. Bartal, A. Fiat, H. Karloff and R. Vohra, “New Algorithms for an Ancient Scheduling Problem,” *Proc. 24th ACM Symposium on the Theory of Computing*, Victoria, Canada, 1992, 51-58.
- [2] U. Faigle, W. Kern and György Turán, “On the Performance of On-Line Algorithms for Partition Problems,” *Acta Cybernetica* 9 (1989), 107-119.
- [3] R. L. Graham, “Bounds for Certain Multiprocessing Anomalies,” *Bell System Technical Journal* 45 (1966), 1563-1581.
- [4] R. L. Graham, E. L. Lawler, J. K. Lenstra, and A. H. G. Rinnooy Kan, “Optimization and Approximation in Deterministic Sequencing and Scheduling: a Survey,” *Annals of Discrete Mathematics* 5 (1979), 287-326.
- [5] D. Karger, S. Phillips, and E. Torng, “A Better Algorithm for an Ancient Scheduling Problem,” manuscript, Stanford University.
- [6] E. L. Lawler, “Recent Results in the Theory of Machine Scheduling,” in A. Bachem, M. Grotchel, and B. Korte (eds.), *Math Programming: The State of the Art (Bonn 1982)*, Springer-Verlag, New York, 1983, 202-234.
- [7] E. L. Lawler, J. K. Lenstra, A. H. G. Rinnooy Kan, and D. B. Shmoys, “Sequencing and Scheduling: Algorithms and Complexity,” to appear in *Handbook of Operations Research and Management Science, Volume IV: Production Planning and Inventory*, S. C. Graves, A. H. G. Rinnooy Kan, and P. Zipkin (eds.), North-Holland.
- [8] J. K. Lenstra and A. H. G. Rinnooy Kan, “An Introduction to Multiprocessor Scheduling,” Technical Report, CWI, Amsterdam, 1988.